

# Breast Cancer Diagnosis

Mar 20, 2022

## Data import

```
# --- Import data ---
breast_cancer =
  read_csv("./data/breast-cancer.csv") %>%
  dplyr::select(-c(1, 33)) %>%
  janitor::clean_names() %>%
  # add extra row
  add_row(diagnosis = 'B', radius_mean = 7.76, texture_mean = 24.54,
    perimeter_mean = 47.92, area_mean = 181, smoothness_mean = 0.05263,
    compactness_mean = 0.04362, concavity_mean = 0,
    concave_points_mean = 0, symmetry_mean = 0.1587,
    fractal_dimension_mean = 0.05884, radius_se = 0.3857,
    texture_se = 1.428, perimeter_se = 2.548, area_se = 19.15,
    smoothness_se = 0.007189, compactness_se = 0.00466, concavity_se = 0,
    concave_points_se = 0, symmetry_se = 0.02676,
    fractal_dimension_se = 0.002783, radius_worst = 9.456,
    texture_worst = 30.37, perimeter_worst = 59.16, area_worst = 268.6,
    smoothness_worst = 0.08996, compactness_worst = 0.06444,
    concavity_worst = 0, concave_points_worst = 0,
    symmetry_worst = 0.2871, fractal_dimension_worst = 0.07039) %>%
  mutate(diagnosis =
    as.numeric(as.factor(recode(diagnosis, `M` = 1, `B` = 0))) - 1) %>%
  mutate_each_(funs(scale(.)), c(2:31))

# --- Data partition (8:2) ---
set.seed(2022)
trRows <- createDataPartition(breast_cancer$diagnosis, p = 0.8, list = FALSE)

# --- `mean` predictors for Newton-Raphson method ---
pred_1 <- as.tibble(breast_cancer[2:11])
bc_scale1 <- as.matrix(cbind(intercept = rep(1, nrow(breast_cancer)),
  pred_1, outcome = breast_cancer$diagnosis))

# train
bc_scale1 <- bc_scale1[trRows, ]
bc_scale1_x <- bc_scale1[, -12]
bc_scale1_y <- bc_scale1[, 12]

# test
t_bc_scale1_x <- bc_scale1[-trRows, -12]
t_bc_scale1_y <- bc_scale1[-trRows, 12]
```

```

# --- All predictors for lasso & CV ---
pred_2 <- as.tibble(breast_cancer[2:31])
bc_scale2 <- as.matrix(cbind(intercept = rep(1, nrow(breast_cancer)),
                             pred_2, outcome = breast_cancer$diagnosis))

# train
bc_scale2 <- bc_scale2[trRows, ]
bc_scale2_x <- bc_scale2[, -32]
bc_scale2_y <- bc_scale2[, 32]

# test
t_bc_scale2_x <- bc_scale2[-trRows, -32]
t_bc_scale2_y <- bc_scale2[-trRows, 32]

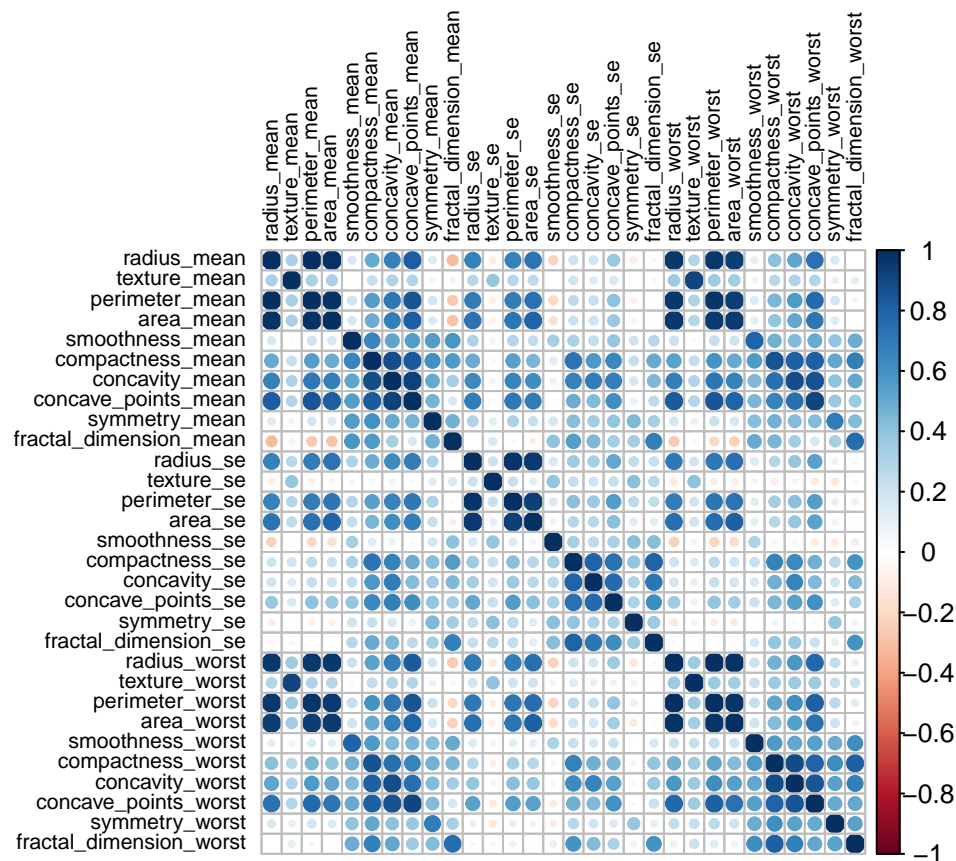
```

## Exploratory Data Analysis

```

# --- Correlation plot ---
corrplot::corrplot(cor(breast_cancer[2:31]), type = "full",
                    tl.cex = .7, tl.col = "black")

```

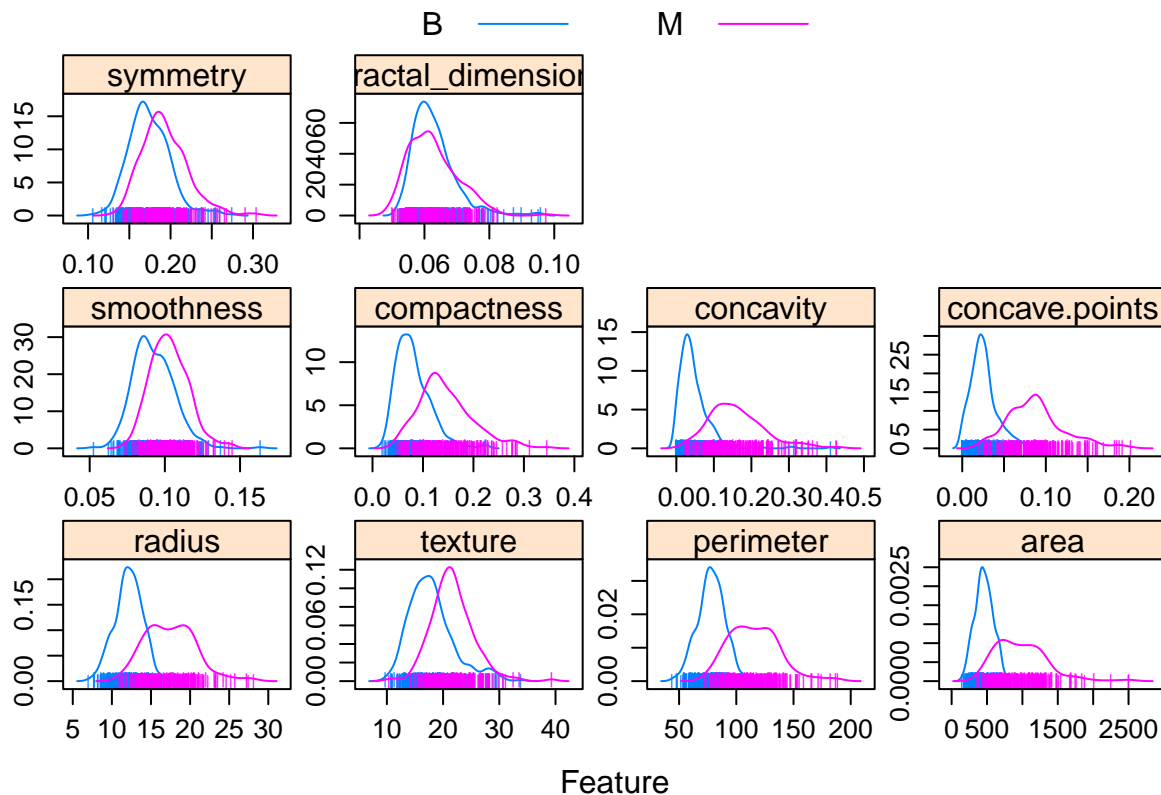


```

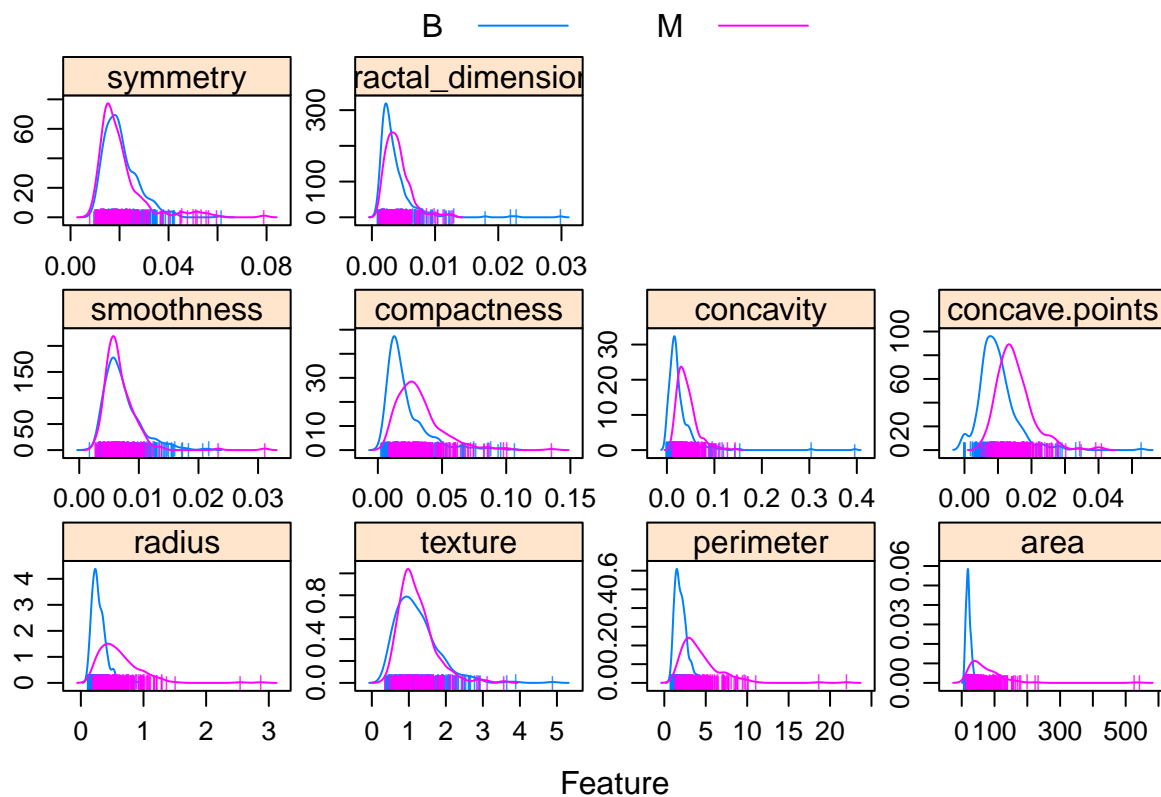
# --- Fea
bc = read.csv("./data/breast-cancer.csv") %>%
  mutate(diagnosis = factor(diagnosis))

```

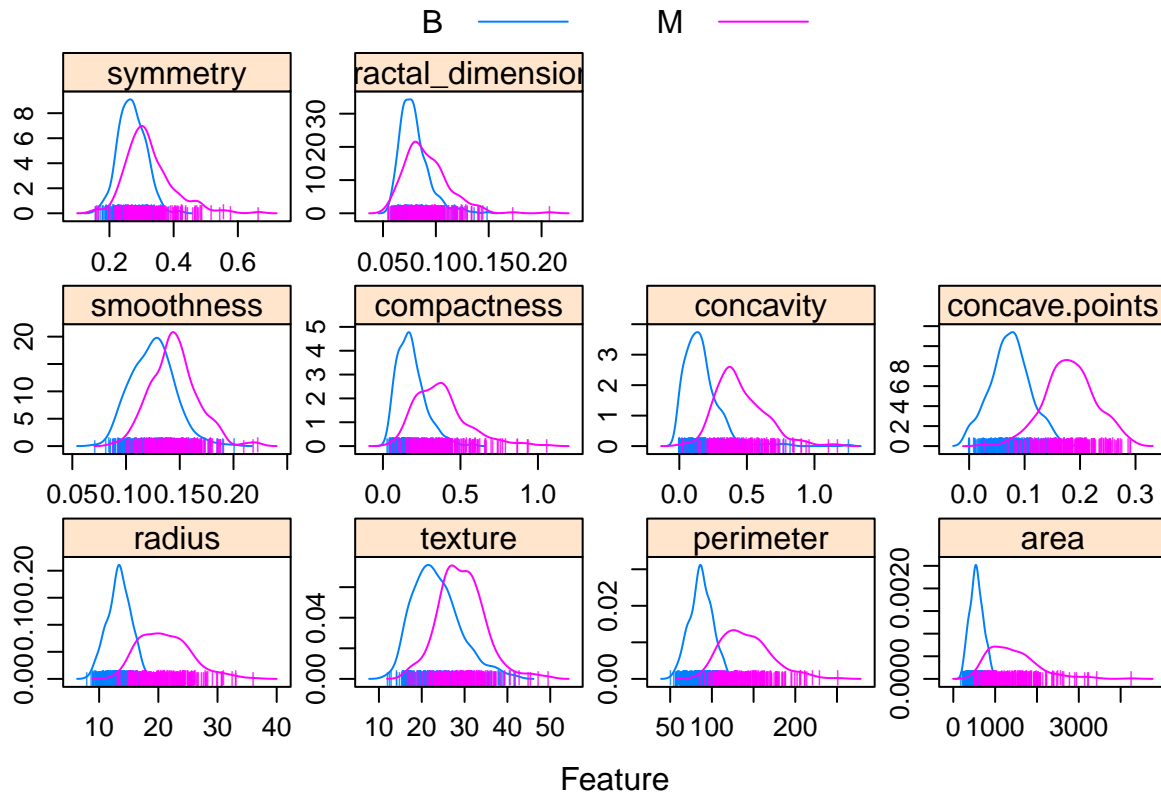
```
#mean
cancer_mean = bc[, 3:12] %>% as_tibble()
colnames(cancer_mean) = gsub("_mean", "", colnames(cancer_mean))
featurePlot(x = cancer_mean,
            y = bc$diagnosis,
            scales = list(x = list(relation = "free"),
                          y = list(relation = "free")),
            plot = "density", pch = "|",
            auto.key = list(columns = 2))
```



```
#se
cancer_se = bc[, 13:22] %>% as_tibble()
colnames(cancer_se) = gsub("_se", "", colnames(cancer_se))
featurePlot(x = cancer_se,
            y = bc$diagnosis,
            scales = list(x = list(relation = "free"),
                          y = list(relation = "free")),
            plot = "density", pch = "|",
            auto.key = list(columns = 2))
```



```
#largest
cancer_max = bc[, 23:32] %>% as_tibble()
colnames(cancer_max) = gsub("_worst", "", colnames(cancer_max))
featurePlot(x = cancer_max,
            y = bc$diagnosis,
            scales = list(x = list(relation = "free"),
                          y = list(relation = "free")),
            plot = "density", pch = "|",
            auto.key = list(columns = 2))
```



## Modified Newton-Raphson method

```
library(MASS)
# Compute likelihood function, gradient, and Hessian matrix
compute_stat <- function(dat, betavec){
  x <- dat # include intercept
  y <- bc_scale1_y
  u <- x %*% betavec # n * 1 vector
  pi <- exp(u) / (1 + exp(u))
  # Log-likelihood
  loglik <- sum(y * u - log(1 + exp(u)))
  # Gradient -- (p + 1) * 1 vector
  grad <- t(x) %*% (y - pi)
  # Hessian -- (p + 1) * (p + 1) matrix
  hess <- -t(x) %*% diag(c(pi * (1 - pi))) %*% x

  return(list(loglik = loglik, grad = grad, hess = hess))
}

# Modified Newton-Raphson method
NewtonRaphson <- function(dat, func, start, tol = 1e-15, maxiter = 200) {
  i <- 0 # Iteration indicator
  cur <- start # Current position / beta's

  # Compute log-likelihood, gradient, and Hessian matrix
  stuff <- func(dat, cur)
```

```

res <- c(0, stuff$loglik, cur) # Store results
prevloglik <- -Inf # Ensure iterations

while (i < maxier && abs(stuff$loglik - prevloglik) > tol) {
  # --- Base case ---
  i <- i + 1
  step <- 1 # Original step size
  prevloglik <- stuff$loglik
  prev <- cur
  cur <- prev - solve(stuff$hess) %*% stuff$grad

  # Step halving and re-direction
  while (func(dat, cur)$loglik < stuff$loglik) {
    # Check if all eigenvalues are negative
    # IF NOT, redirection
    eig <- eigen(stuff$hess)$value
    if (sum(eig < 0) != length(eig)) {
      gamma <- max(eig)
      new_hess <- stuff$hess - gamma * diag(nrow = dim(dat)[1])
      cur <- prev - solve(new_hess) %*% stuff$grad
    }
    # Reduce step size by half
    else {
      step <- step / 2
      cur <- prev - step * solve(stuff$hess) %*% stuff$grad
    }
  }
  stuff <- func(dat, cur)

  # break if NA
  if (is.na(stuff$loglik)) {break}

  res <- rbind(res, c(i, stuff$loglik, cur))
}
return(res)
}

```

NR results (both 10 means & all 30 predictors)

```

# NewtonRaphson(bc_scale1_x, compute_stat, rep(1, 11))
# NewtonRaphson(bc_scale2_x, compute_stat, rep(10, 31))

# The result of using 30 of predictors is not ideal
# due to some potential collinearity

```

Compare with glm results

```

# Fit a logistic model
glm_fit <- glm(diagnosis ~., family = binomial(link = "logit"), data = breast_cancer[, 1:11])
glm_fit$coefficients

```

```
##          (Intercept)          radius_mean          texture_mean
##          0.48701675          -7.22185053          1.65475615
##          perimeter_mean          area_mean          smoothness_mean
##          -1.73763027          14.00484560          1.07495329
##          compactness_mean          concavity_mean          concave_points_mean
##          -0.07723455          0.67512313          2.59287426
##          symmetry_mean fractal_dimension_mean
##          0.44625631          -0.48248420
```

```
# Comparison
```

```
NR_result <- NewtonRaphson(bc_scale1_x, compute_stat, rep(10, 11))
round(tail(NR_result, 1)[, 3:13], 3) == round(as.vector(glm_fit$coefficients), 3) # return all TRUE
```

```
## [1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

## Soft threshold and coor-wise Lasso

```
# Soft Threshold function
```

```
sf <- function(beta, lambda) {
  beta <- ifelse(lambda < abs(beta),
    ifelse(beta > 0, beta - lambda, beta + lambda), 0)
  return(beta)
}
```

```
# Coordinate-wise LASSO with fixed lambda
```

```
y <- breast_cancer$diagnosis

cd_lasso <- function(dat, y, betavec, lambda, maxier = 2000, tol = 1e-10) {
  x <- dat # n * (p + 1) matrix, standardized
  i <- 0
  objfun <- 0
  prev_objfun <- -Inf # ensure iterations
  res <- c(0, objfun, betavec)

  while (i < maxier && abs(objfun - prev_objfun) > tol) {
    i <- i + 1
    prev_objfun <- objfun

    for (j in 1:length(betavec)) {
      u <- x %*% betavec
      pi <- exp(u) / (1 + exp(u))
      # working weights
      w <- pi * (1 - pi)
      w <- ifelse(w < 1e-5, 1e-5, w)
      # working response
      z <- x %*% betavec + (y - pi) / w
      z_dej <- x[, -j] %*% betavec[-j]
      betavec[j] <-
        sf(sum(w * x[, j] * (z - z_dej)), lambda) / (sum(w * x[, j]^2))
    }
  }
}
```

```

objfun <-
  sum(w * (z - x %*% betavec)^2) / (2 * dim(x)[1]) + lambda * sum(abs(betavec))
  if (is.na(objfun)) {break}

  res <- rbind(res, c(i, objfun, betavec))
}

return(res)
}

```

Try `cd_lasso...`

```

# cd_lasso(bc_scale2_x, bc_scale2_y, betavec = rep(1, 31), lambda = 1)

```

## path-wise coordinate-wise for lasso solution path

```

# pathwise based on cd_lasso
pathwise_cd_lasso <- function(dat, y, betavec, lambda) {

  sort(lambda, decreasing = TRUE)
  pathwise_res <- NA

  for (j in 1:length(lambda)) {
    cd_results <- cd_lasso(dat, y, betavec, lambda[j])

    # warm start
    min_objfun_ind <- which.min(cd_results[-1, 2]) + 1
    betavec <- cd_results[min_objfun_ind, -c(1, 2)]
    objfun <- cd_results[min_objfun_ind, 2]

    pathwise_res <- rbind(pathwise_res,
                          c(j, lambda = lambda[j], obj = objfun,
                            beta = betavec))
  }

  return(pathwise_res[-1, ])
}

```

## 5-fold CV using `cd_lasso`

```

cv = function(raw_data, lambda_vec, kfolds = 5, betavec){

  fold_index = cut(seq(1,nrow(raw_data)),breaks = kfolds, labels = FALSE)
  raw_data = raw_data[sample(nrow(raw_data)), ]

  final_res <- c(lambda = 0, auc = 0, rss = 0)

```



```

for (i in 1:length(lambda_vec)) {
  curlambda = lambda_vec[i]
  fold_auc = vector(); fold_rss = vector()

  for (k in 1:kfolds) {
    train = raw_data[fold_index != k, ]
    validation = raw_data[fold_index == k, ]

    train_x = train[, -ncol(raw_data)]
    train_y = train[, ncol(raw_data)]
    validation_x = validation[, -ncol(raw_data)]
    validation_y = validation[, ncol(raw_data)]

    # cd_lasso
    res = cd_lasso(train_x, train_y, betavec, curlambda)

    curbeta = res[dim(res)[1], 3:dim(res)[2]]
    u = validation_x %*% as.matrix(curbeta) # u: n * 1 vector
    prob = exp(u) / (1 + exp(u))

    fold_auc[k] = roc(as.factor(validation_y), prob)$auc[1]
    fold_rss[k] = sum((validation_y - prob)^2)
  }

  final_res = rbind(final_res, c(lambda = curlambda, mean_auc = mean(fold_auc),
                                sd_auc = sd(fold_auc), rss = mean(fold_rss)))
}

return(final_res)
}

set.seed(2022)
cv_res <- cv(bc_scale2, lambda_vec = exp(seq(4, -3, length = 30)), betavec = rep(1, 31))

```

Compare optimal with full using test data

```

opt_lambda = cv_res[which.max(cv_res[, 2]), 1]

# Optimal lambda
optimal = cd_lasso(t_bc_scale2_x, t_bc_scale2_y, lambda = opt_lambda, betavec = rep(1,31))[,-c(1:2)] %>%
  as.matrix()
u_opt <- t_bc_scale2_x %*% optimal[nrow(optimal),]
prob_opt <- as.matrix(exp(u_opt) / (1 + exp(u_opt)))
roc(as.factor(t_bc_scale2_y), prob_opt)$auc[1]

## [1] 0.9988981

# Without lambda
full = cd_lasso(t_bc_scale2_x, t_bc_scale2_y, lambda = 0, betavec = rep(1,31))[,-c(1:2)] %>%
  as.matrix()

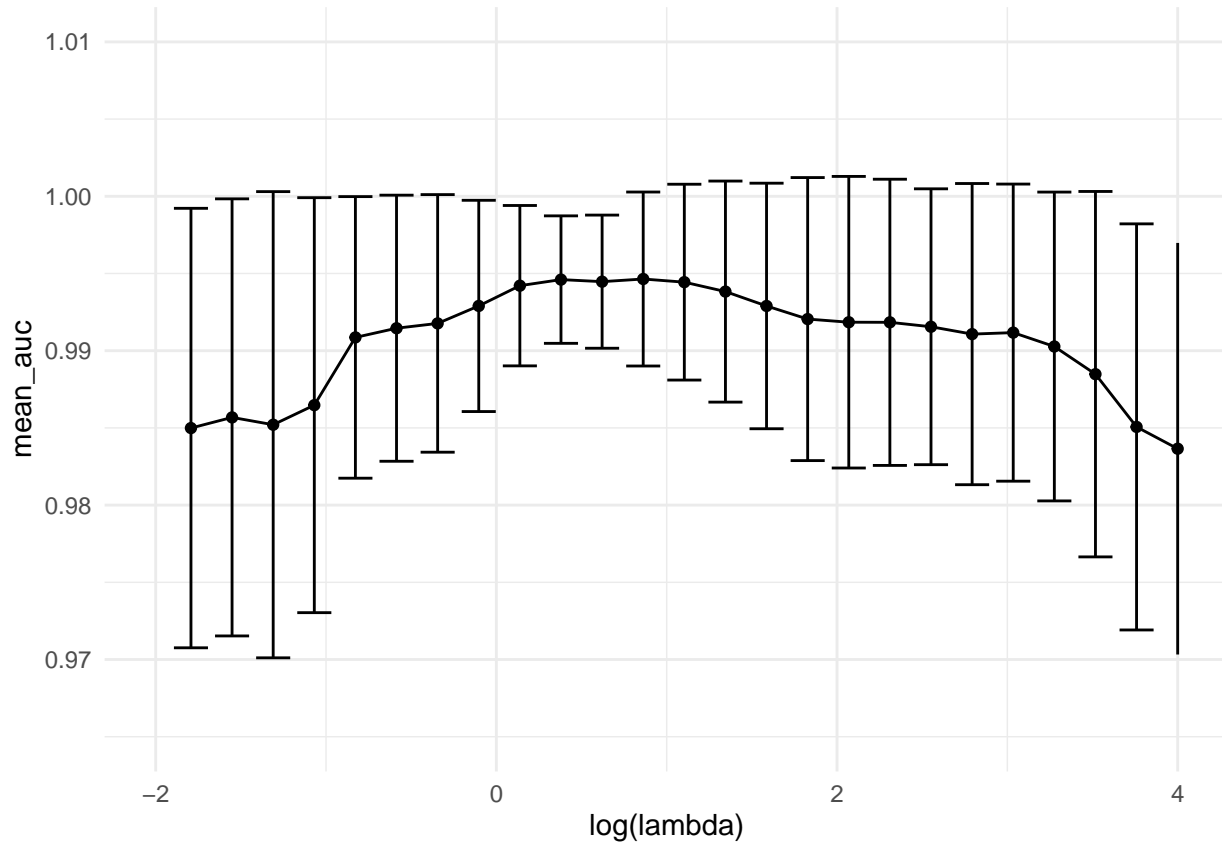
```

```
u_full <- t_bc_scale2_x %*% full[nrow(full),]
prob_full <- as.matrix(exp(u_full) / (1 + exp(u_full)))
roc(as.factor(t_bc_scale2_y), prob_full)$auc[1]
```

```
## [1] 1
```

5-fold cv mean auc & sd auc

```
# mean & sd
as.tibble(cv_res[-1, ]) %>%
  ggplot(aes(x = log(lambda), y = mean_auc)) +
  geom_point() +
  geom_line() + xlim(c(-2, 4)) +
  geom_errorbar(aes(ymin = mean_auc - sd_auc, ymax = mean_auc + sd_auc), width = .2) +
  ylim(c(0.965, 1.01)) + theme_minimal()
```



Lasso solution path

```
# Compute maximum lambda shrinking every beta to 0
max_lambda2 <- max(t(bc_scale2_x) %*% bc_scale2_y)

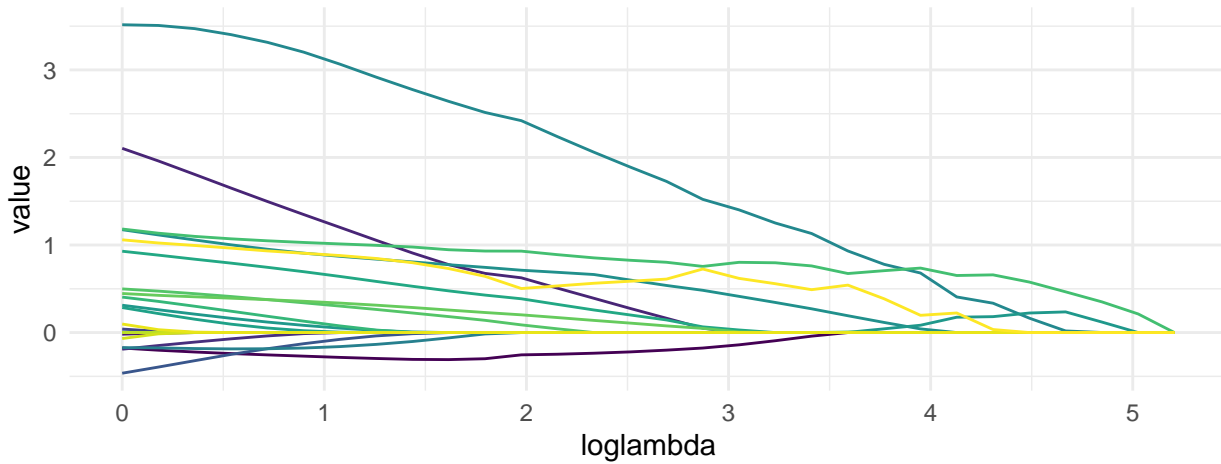
# pathwise cd_lasso
```

```

pathwise = pathwise_cd_lasso(bc_scale2_x, bc_scale2_y, rep(1,31), exp(seq(log(max_lambda2), 0, length =
tidy.path = as.tibble(pathwise)[,-c(1,3)] %>%
  pivot_longer(2:32, names_to = "beta", values_to = "value") %>%
  mutate(loglambda = log(lambda))

ggplot(data = tidy.path, aes(x = loglambda, y = value, color = beta, group = beta)) +
  geom_line()

```



beta

— beta1	— beta16	— beta22	— beta29	— beta7
— beta10	— beta17	— beta23	— beta3	— beta8
— beta11	— beta18	— beta24	— beta30	— beta9
— beta12	— beta19	— beta25	— beta31	
— beta13	— beta2	— beta26	— beta4	
— beta14	— beta20	— beta27	— beta5	
— beta15	— beta21	— beta28	— beta6	